

What is pre-emptive scheduling?

What is a dispatch loop?

What are the three levels of parallelism granularity?

How can they be further categorized?

What are some motivations of concurrency?

What is a race condition?

What is the purpose of synchronization?

What is the memory coherence problem?
Who does it affect?

What are the two broad categories of single-site (non-distributed) parallel computers?

Define *task*.

Contrast heavyweight and lightweight threads.